

Appendix

The material presented in this appendix is meant to provide additional details on the construction of both room similarity networks and co-occurrence networks. Both methods and the justification for certain choices are presented here.

Room Similarity Networks

For constructing networks of room similarity we use the Brainerd-Robinson (B-R) coefficient to define the scale of similarity among rooms where the similarity (S) between room a and room b is defined as follows:

$$S_{ab} = \frac{200 - \sum_k |P_{ak} - P_{bk}|}{200}$$

where k is all artifact classes, P_{ak} is the percent of class k in room a , and P_{bk} is the percent of class k in room b (Brainerd 1951; Robinson 1951). Specifically, we converted the counts of the 13 artifact classes from each room to percentages and then calculated the Brainerd-Robinson (B-R) similarity coefficient (Brainerd 1951; Robinson 1951) among every possible pairwise combination of rooms. The result is a symmetric matrix of B-R similarity scores among all rooms, with the number of rows and columns equal to the number of rooms.

This matrix was then converted to a weighted network graph with the weight of ties between nodes (rooms) equal to the B-R similarity of that room pair. The network was then binarized by removing all ties with values below 75% similarity (a B-R score below 150). This threshold of 75% is arbitrary, but was selected based on examination of histograms of overall similarity scores across the entire dataset in order to capture the strongest (most similar) relationships among rooms. While it is often important to use weighted ties (which indicate the strength or value of the relationship between two nodes) in network analysis of archaeological

materials, rather than ties based on presence-absence data (see Peebles and Roberts 2013), the current research is concerned primarily with the strongest associations among different artifact classes. Weak ties are less useful in evaluating these relationships since they are more likely to be driven by sampling error (especially when considering rare objects with low overall frequencies) and associations among the most common objects. Consequently, use of binarized data based on excluding weaker ties (below 75 percent similarity in this case) provides both useful visuals and a powerful tool for reducing complexity in the overall dataset by highlighting only the strongest similarities and associations (see also discussions in Hill et al. 2015 and Weidele et al. 2016).

One consequence of examining only these strong ties (greater than 75% similarity) is that the resulting network graphs tend to include many small clusters of rooms or individual rooms isolated from the rest of the overall network. These clusters and individual rooms would otherwise have been connected to the rest of the network through a number of weakly similar relationships, but become disconnected when examining only strong ties. In this case, the use of only strong ties helps us to identify the most internally cohesive groupings of rooms in the network. Since the ties in this analysis are defined on the basis of a similarity measure (the B-R coefficient), internal cohesion of a group indicates high similarity of all rooms in that group. This exclusion of rooms connected to the network only by weak ties is therefore a desirable outcome in evaluating strongly consistent patterning of archaeological material, rather than focusing on outliers and weaker patterns. An examination of the rooms with no strong ties to other rooms suggests that their isolation is largely driven by high proportions of the rarest artifact classes. While these isolated rooms may offer interesting insights into the diversity of depositional practices within the site, they do not help to characterize the primary depositional behaviors that

resulted in the site assemblage and which is the primary focus of the current research.

Clustering within the room similarity networks was evaluated by applying the Girvan-Newman method to each network graph. This method calculates the centrality of each tie in the network and breaks the network into clusters by removing the most central relationships. This process effectively acts like traditional cluster analysis except working on a subtractive, rather than additive, principle. The primary benefit of the Girvan-Newman analysis in comparison with cluster analysis is that, as a subtractive analysis, it is more likely to accurately assign peripheral members to a cluster. As an additive analysis, cluster analysis is more likely to miss-assign these peripheral members of a cluster (Newman and Girvan 2004:2). The use of the Girvan-Newman analysis follows the same justification as our use of a 75% threshold in constructing the network, as a method for isolating the most internally cohesive clusters of rooms and strongest patterns of relationships within the network.

Importantly, as in any cluster analysis, the number of clusters produced by the Girvan-Newman technique is determined by the researcher through the selection of a modularity value, producing anywhere from a minimum number of clusters (of one) up to as many clusters as there are nodes in the network. However, a modularity score can be calculated to evaluate how much variance is accounted for by the addition of each cluster. This modularity score is based on eigenvalues calculated for each cluster added to the network (Newman 2006). For instance, the NBNK dataset the maximum modularity score was achieved for a seven-cluster network (**Figure 2** in the main text), accounting for 67.90 percent of the variance in the data.

Co-Occurrence Measurement

The measure of co-occurrence used in constructing our co-occurrence networks was calculated using a binomial model defined for artifact classes *a* and *b* as:

$$C_{ab} = \frac{o - Np_{ab}}{\sqrt{Np_{ab}(1 - p_{ab})}}$$

71 o = the observed number of co-occurrences between artifact classes a and b

72 N = the total number of assemblages

73 p = the expected proportion of co-occurrences for artifact classes a and b defined as the
 74 expected proportion of occurrences for class a multiplied by the expected proportion of
 75 occurrences for artifact class b

76

77 Observed values of co-occurrence can be obtained by multiplying the binary version of the data
 78 matrix by its transpose. The expected proportion of occurrence for each artifact class is defined
 79 simply as the total number of occurrences divided by the number of assemblages. Thus, Np_{ab}
 80 represents the expected count of co-occurrences between class a and class b .

81 This measure, originally described by Kintigh (2006), provides an index of the number of
 82 co-occurrences observed in relation to the number that would be expected by chance given the
 83 overall frequency of each artifact class. This measure approximates a Z-score such that the
 84 absolute values can be interpreted in standard deviation units above and below 0. Thus, we might
 85 expect the distribution of values for a given table of occurrence data to have a mean of
 86 approximately 0 with roughly equivalent numbers of positive and negative values. As **Table 2** in
 87 the main text of this article illustrates, however, the range of values obtained from the NBNK
 88 dataset was highly positively skewed (skewness = 1.36; see also **Figure S.1**). The values
 89 obtained for the other datasets evaluated in this research were similarly skewed. In order to
 90 assess the potential factors driving this skewed distribution, we conducted a simulation
 91 experiment.

92 First, we created 10,000 randomized versions of the original $n \times m$ matrix of occurrences
 93 (n = number of rooms and m = number of artifact classes) for the NBNK dataset by permuting
 94 each column independently. Thus, this procedure resulted in randomized co-associations while
 95 preserving the relative frequencies of artifact classes. Next, we calculated our measure of co-


occurrence C as described above for each random matrix, resulting in 10,000 $m \times m$ matrices of co-occurrence scores. The distribution of values obtained for all of these random runs combined is normally distributed (**Figure S.2**) with a mean of 0. In only 0.42% of all randomized matrices did we obtain a skewness greater than or equal to the observed value of 1.36 in our original matrix. Importantly, all random runs that resulted in highly positively skewed distributions (>1) were characterized by multiple co-associations among the 5 rarest artifact classes including spindle whorls, crystals, jet, copper, and cylinder jars whereas such co-associations were less common in random runs with skewness values closer to 0. Indeed, **Figure S.3** shows the average skewness value increases in relation to the number of co-occurrences of the rarest artifact classes. If we return to our original dataset we find a total of 10 co-occurrences among the 5 rarest artifact classes (higher than almost all random runs). Overall, this suggests that the positive skew of C values in our original observed table is driven by the tendency for rare objects to cluster in the NBNK dataset and is not simply an artifact of the method of assessing co-occurrence. The result of this procedure is a table with artifact classes as both rows and columns and with each cell representing the co-occurrence index between two artifact classes. See **Table 2** for an example from the NBNK dataset.

Network Cliques

After converting the matrix of co-occurrence values into a network graph, as described in the main text, clustering within the network was evaluated using the network concept of a “clique”. A clique in network terms is a “maximal, complete subgraph of three or more nodes” (Wasserman and Faust 1994:254). In this case, maximal means that no additional nodes can be added to the clique that meet the criteria of the clique (having ties to every other clique member). Due to their definition, two or more cliques can share members (though no clique can ever be

entirely contained within another, given the “maximal” definition of a clique). Since our analysis is concerned with behavioral sets of things, overlap between cliques is a non-issue given that use of objects from an artifact class in one context does not inherently preclude use of that same artifact class in a different context. The clique concept is useful for an assemblage-oriented approach because it moves beyond the evaluation of single artifact classes or pairings of artifact classes to look at entire sets of co-occurrence within an assemblage.

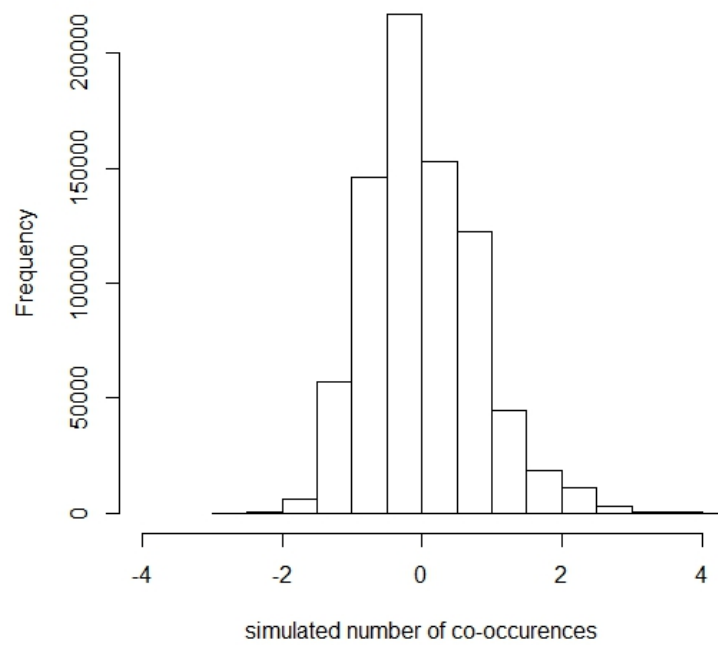
In weighted networks the clique can be expanded to evaluate cliques hierarchically based on the values of ties in a clique. In this case the clique definition is modified such that a ‘clique at level C’ meets not only the basic definition of a clique, but also requires that all ties between members of the clique ‘at level C’ are of value C or greater (Wasserman and Faust 1994: 278). Applied to the network of co-occurrence indices, a clique in the network is a set of three-or-more artifact classes which all co-occur with each-other more frequently than expected by chance at a given threshold. The use of valued or weighted ties then helps identify the most cohesive cliques, which for the co-occurrence network are only those cliques which have high values of co-occurrence between all clique members.



148

149

Figure S.1. Distribution of C co-occurrence values for original data table.



150

151

152

Figure S.2. Distribution of C co-occurrence values for the 10,000 randomized matrices.




Figure S.3. Plot showing the average skewness and confidence intervals for 10,000 randomized datasets in relation to the number of co-occurrences among the 5 rarest artifact classes.


```

175
176
177 R function for calculating C index:
178 Input – Matrix x of n contexts/rooms as rows and m artifact classes as columns.
179 Data are assumed to be 0/1 presence/absence.
180
181 cooccur <- function(x) {
182
183   # calculate the proportional occurrence of each artifact class
184   nm.p <- colSums(x)/nrow(x)
185
186   # calculated observed co-occurrences through matrix multiplication
187   obs <- t(as.matrix(x)) %*% (as.matrix(x))
188   diag(obs) <- 0
189   # create matrix of expected values based on proportional occurrence
190   expect <- matrix(0,nrow(obs),ncol(obs))
191   for (i in 1:nrow(obs)) {
192     for (j in 1:ncol(obs)) {
193       expect[i,j] <- (nm.p[i]*nm.p[j])*nrow(x)}
194
195   # convert expected count to expected proportion
196   p <- expect/nrow(x)
197
198   #calculate the denominator of the equation
199   denm <- sqrt(expect*(1-p))
200
201   # calculate final matrix of scores and output
202   out <- (obs-expect)/denm
203   diag(out) <- 0
204   return(out)}
205
206 ## end code
207

```